

```

1. import torch # pytorch main library
2. import torchvision # computer vision utilities
3. import torchvision.transforms as transforms # transforms used in the pre-processing of
   the data
4.
5. import torch.nn as nn
6. import torch.nn.functional as F
7.
8. import matplotlib.pyplot as plt
9. import numpy as np
10.
11. import torch.optim as optim
12.
13. # Check if GPU is available
14. device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
15.
16. # Assuming that we are on a CUDA machine, this should print a CUDA device:
17. print(device)

```

This code imports several libraries, including PyTorch (torch), a library for building machine learning models, and torchvision, a library containing computer vision utilities. The torchvision.transforms module is also imported, which contains preprocessing functions for image data. The code also imports several PyTorch modules, including torch.nn, which contains classes for building neural networks, and torch.nn.functional, which contains functions for building neural network layers. The code also imports the matplotlib and numpy libraries for plotting and numerical operations, respectively.

The code then checks if a GPU (Graphics Processing Unit) is available for use. If a GPU is available, the device variable is set to 'cuda:0', which tells PyTorch to run computations on the GPU. If a GPU is not available, the device variable is set to 'cpu', which tells PyTorch to run computations on the CPU.

Finally, the code prints the value of the device variable, which should indicate whether the computations will be run on the GPU or the CPU.

```

1. # Function to get the statistics of a dataset
2. def get_dataset_stats(data_loader):
3.     mean = 0.
4.     std = 0.
5.     nb_samples = 0.
6.     for data in data_loader:
7.         data = data[0] # Get the images to compute the statistics
8.         batch_samples = data.size(0)
9.         data = data.view(batch_samples, data.size(1), -1)
10.        mean += data.mean(2).sum(0)
11.        std += data.std(2).sum(0)
12.        nb_samples += batch_samples
13.
14.    mean /= nb_samples
15.    std /= nb_samples
16.    return mean, std
17.

```

```

18. # functions to show an image
19. def imshow(img, stats):
20.     img = img * stats[1] + stats[0]    # unnormalize
21.     npimg = img.numpy() # convert the tensor back to numpy
22.     plt.imshow(np.transpose(npimg, (1, 2, 0)))
23.     plt.show()

```

This code defines two functions: `get_dataset_stats` and `imshow`.

The `get_dataset_stats` function takes a single argument, `data_loader`, which is a PyTorch  `DataLoader` object that iterates over a dataset and returns a batch of data. The function initializes three variables: `mean`, `std` and `nb_samples`. The `mean` and `std` variables will be used to store the mean and standard deviation of the dataset's data, respectively. The `nb_samples` variable will be used to count the total number of samples in the dataset.

The function then iterates over the data in the `data_loader`, which is a batch of data, and for each iteration, it does the following:

1. It extracts the images from the batch of data, which is stored in the variable `data`.
2. It gets the number of samples in the batch, stored in the variable `batch_samples`.
3. It reshapes the data from 3D tensor to 2D tensor by using the `view` method and flattening the last dimension.
4. It accumulates the mean of each channel over all samples in the batch.
5. It accumulates the standard deviation of each channel over all samples in the batch.
6. It increments the total number of samples by the number of samples in the current batch.

After all the batches have been processed, the function calculates the mean and standard deviation by dividing the accumulated mean and standard deviation by the total number of samples, respectively. Finally, the function returns the mean and standard deviation of the dataset's data.

The `imshow` function takes two arguments, `img` and `stats`. The `img` argument is a PyTorch tensor containing an image, and the `stats` argument is a tuple of two elements, which are the mean and standard deviation of the dataset's data, respectively. The function does the following:

1. It un-normalizes the image by multiplying it with the standard deviation and adding the mean.
2. It converts the tensor back to a numpy array using the `numpy()` method.
3. It transposes the numpy array so that the channels are in the last dimension.
4. It shows the image using the `plt.imshow()` function from the `matplotlib` library.
5. It shows the plot using the `plt.show()` function from the `matplotlib` library.

The `imshow` function allows to display the image in its original format, thus the mean and std are used to un-normalize the image before displaying it.

```
1. batch_size = 256
2.
3. transform = transforms.Compose(
4.     [transforms.ToTensor()]) # Convert the data to a PyTorch tensor
5.
6. # Load development dataset
7. devset = torchvision.datasets.MNIST(root='./data', train=True,
8.                                     download=True, transform = transform)
9.
10. train_set_size = int(len(devset) * 0.8)
11. val_set_size = len(devset) - train_set_size
12.
13. # Split the development set into train and validation
14. trainset, valset = torch.utils.data.random_split(devset, [train_set_size,
15.    val_set_size], generator=torch.Generator().manual_seed(42))
16.
17. # Get the data loader for the train set
18. trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
19.    shuffle=True, num_workers=2)
20.
21. # Compute the statistics of the train set
22. stats = get_dataset_stats(trainloader)
23. print("Train stats:", stats)
24. # Pre-processing transforms
25. transform = transforms.Compose(
26.     [transforms.ToTensor(),
27.      transforms.Normalize((stats[0]), (stats[1]))])
28.
29.
30. # Load the development set again using the proper pre-processing transforms
31. devset = torchvision.datasets.MNIST(root='./data', train=True,
32.                                     download=True, transform = transform)
33.
34. # Split the development set into train and validation
35. trainset, valset = torch.utils.data.random_split(devset, [train_set_size,
36.    val_set_size], generator=torch.Generator().manual_seed(42))
37.
38. # Get the data loader for the train set
39. trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
40.    shuffle=True, num_workers=2)
41.
42. # Get the data loader for the test set
43. valloader = torch.utils.data.DataLoader(valset, batch_size=batch_size,
44.    shuffle=True, num_workers=2)
45.
46. # Get the test set
47. testset = torchvision.datasets.MNIST(root='./data', train=False,
48.    download=True, transform=transform)
49.
50. # Get the data loader for the test set
51. testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
52.    shuffle=False, num_workers=2)
53.
54. classes = ('zero', 'one', 'two', 'three',
55.            'four', 'five', 'six', 'seven', 'eight', 'nine')
```

This code is for loading, splitting, normalizing and pre-processing the MNIST dataset for training, validation and testing a model. It starts by defining the batch size as 256, which means that the dataset will be divided into batches of 256 samples each.

Then, it defines a transform using the `torchvision.transforms.Compose` function, which applies multiple pre-processing steps to the data. In this case, it applies the `torchvision.transforms.ToTensor()` function, which converts the input data from a PIL image to a PyTorch tensor.

Next, it loads the MNIST dataset, which is a dataset of handwritten digits, using the `torchvision.datasets.MNIST` function. The `root` argument is set to  `'./data'`, which is the directory where the dataset will be downloaded and stored. The `train` argument is set to `True`, which means that the dataset will be loaded as a training dataset. The `download` argument is set to `True`, which means that the dataset will be downloaded if it is not already present in the specified directory. The `transform` argument is set to the transform defined earlier.

It then splits the dataset into a train and validation set by using the `torch.utils.data.random_split()` function. The first argument is the dataset, the second argument is a list containing the size of the train set and the validation set, and the third argument is a generator that will use a manual seed of 42 to generate random numbers.

Next, it creates PyTorch `DataLoader` objects for the train, validation and test sets using the `torch.utils.data.DataLoader` function. The first argument is the corresponding dataset, the second argument is the batch size, the third argument is whether to shuffle the data, and the fourth argument is the number of worker threads to use for loading the data.

Then, it computes the statistics of the train set using the `get_dataset_stats` function, which returns a tuple of mean and std. Then, it normalizes the train, validation and test sets using the computed mean and std and the `torchvision.transforms.Normalize()` function, which normalizes the data by subtracting the mean and dividing by the standard deviation.

Finally, it defines a tuple of classes containing the names of the digits (zero, one, two, ..., nine) that the model will be trying to classify.

```
1. # get some random training images
2. dataiter = iter(trainloader)
3. images, labels = next(dataiter)
4.
5. # show images
6. imshow(torchvision.utils.make_grid(images[:8]), stats)
7. # print labels
8. print(' '.join(f'{classes[labels[j]]:5s}' for j in range(8)))
```

This code is used to display some random training images from the MNIST dataset and their corresponding labels.

First, it creates an iterator for the training data using the `iter()` function and the `trainloader` object. Then, it uses the `next()` function to get the next batch of data from the iterator, which is stored in the variables `images` and `labels`.

Then, it uses the `torchvision.utils.make_grid` function to create a grid of images from the batch of images, and the function takes the first 8 images from the batch. It then calls the `imshow` function with the grid of images and the statistics of the dataset, to display the images in their original format.

Finally, it uses the `join()` function to concatenate the labels of the images, which are the classes of the digits, and it prints them. The `f'{classes[labels[j]]:5s}'` is used to format the output as a string of 5 characters and it uses the index of the labels to get the class name from the classes tuple.

```
1. class Net(nn.Module):
2.     def __init__(self):
3.         super().__init__()
4.         self.fc1 = nn.Linear(28*28, 120)
5.         self.fc2 = nn.Linear(120, 84)
6.         self.fc3 = nn.Linear(84, 10)
7.
8.     def forward(self, x):
9.         x = torch.flatten(x, 1) # flatten all dimensions except batch
10.        x = F.relu(self.fc1(x))
11.        x = F.relu(self.fc2(x))
12.        x = self.fc3(x)
13.        return x
14.
15.
16. net = Net()
17. net.to(device)
```

This is a PyTorch implementation of a simple feedforward neural network. The `Net` class is a subclass of PyTorch's `nn.Module` class and it defines the architecture of the network. The `__init__` method of the `Net` class initializes three fully connected layers, `fc1`, `fc2` and `fc3`, using PyTorch's `nn.Linear` module. The input and output sizes of these layers are specified as arguments to the `nn.Linear` module, with the input size of the first layer being `28*28` (the size of a single image in the MNIST dataset) and the output size of the last layer being `10` (the number of classes in the MNIST dataset).

The `forward` method of the `Net` class defines the forward pass of the network. It takes an input tensor `x` and applies a series of operations to it: `x` is first flattened using `torch.flatten` so that it can be passed to the first linear layer, `fc1`. Then, the output of `fc1` is passed through a rectified linear unit (ReLU) activation function using `F.relu`. This process is repeated for the second and third linear layers, `fc2` and `fc3`. Finally, the output of the last linear layer `fc3` is returned.

The last line of the code creates an instance of the `Net` class and assigns it to the variable `net`. The `net.to(device)` line moves the model to the specified device (e.g. `'cuda'` for GPU)

```
1. criterion = nn.CrossEntropyLoss() # Loss function
2. optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9) # Optimizer used for training
```

The code above sets up the training process for the neural network defined earlier.

The first line initializes the loss function, which is used to measure the difference between the predicted and actual output during training. In this case, the `nn.CrossEntropyLoss` function is used, which is a commonly used loss function for classification tasks

The second line initializes the optimizer, which is used to update the network's parameters during training. The `optim.SGD` function is used, which stands for stochastic gradient descent. The first argument passed to `optim.SGD` is the parameters of the network that need to be optimized (in this case, `net.parameters()`), the second argument `lr` is the learning rate, which determines the step size at which the optimizer makes updates to the parameters during training, a smaller learning rate means the optimizer will make smaller updates and the training process will be slower. The third argument `momentum` is a parameter used in the optimizer that helps the optimizer converge faster by averaging the gradients across multiple steps.

This optimizer and loss function combination is a common choice for training neural networks, but other options such as Adam, Adagrad, and RMSprop are also commonly used.

```
1. nepochs = 20
2. PATH = './cifar_net.pth' # Path to save the best model
3.
4. best_loss = 1e+20
5. for epoch in range(nepochs): # loop over the dataset multiple times
6.     # Training Loop
7.     train_loss = 0.0
8.     for i, data in enumerate(trainloader, 0):
9.         # get the inputs; data is a list of [inputs, labels]
10.        inputs, labels = data[0].to(device), data[1].to(device)
11.
12.        # zero the parameter gradients
13.        optimizer.zero_grad()
14.
15.        # forward + backward + optimize
16.        outputs = net(inputs)
17.        loss = criterion(outputs, labels)
18.        loss.backward()
19.        optimizer.step()
20.
21.        train_loss += loss.item()
22.    print(f'{epoch + 1}, train loss: {train_loss / i:.3f}', end = ' ')
23.
24.    val_loss = 0
25.    # since we're not training, we don't need to calculate the gradients for our
    outputs
26.    with torch.no_grad():
27.        for i, data in enumerate(valloader, 0):
28.            # get the inputs; data is a list of [inputs, labels]
29.            inputs, labels = data[0].to(device), data[1].to(device)
30.
31.            outputs = net(inputs)
32.            loss = criterion(outputs, labels)
33.
34.            val_loss += loss.item()
35.    print(f'val loss: {val_loss / i:.3f}')
36.
```

```

37.         # Save best model
38.         if val_loss < best_loss:
39.             print("Saving model")
40.             torch.save(net.state_dict(), PATH)
41.
42. print('Finished Training')

```

This code trains the neural network defined earlier using the training data and the optimizer and loss function initialized earlier. The training process is repeated for a specified number of epochs (in this case, 20).

The outer loop iterates over the number of epochs. In each iteration, the training loss is computed by iterating over the training data using the PyTorch DataLoader, trainloader. For each training sample, the gradients of the network's parameters are set to zero using `optimizer.zero_grad()`, the forward pass is performed using `outputs = net(inputs)`, the loss is calculated using the loss function `criterion(outputs, labels)`, the gradients are computed using `loss.backward()`, and the optimizer updates the parameters using `optimizer.step()`. The training loss is accumulated over the entire training set and printed out at the end of each epoch.

The validation loss is also computed by iterating over the validation data using the PyTorch DataLoader, valloader. The validation loss is accumulated over the entire validation set and printed out at the end of each epoch.

If the validation loss is smaller than the best validation loss seen so far, the current model is saved to the specified file path `./cifar_net.pth` using the `torch.save` function.

After the training is complete, the script prints 'Finished Training' to indicate that the training process is done.

```

1. # Load the best model to be used in the test set
2. net = Net()
3. net.load_state_dict(torch.load(PATH))

```

This code loads the best model that was saved during the training process.

It first creates an instance of the `Net` class, which defines the architecture of the network. Then, the `load_state_dict` method is used to load the saved model's parameters, which were previously saved to the file path specified in the training code `PATH = './cifar_net.pth'`. The `torch.load` function is used to load the saved model's parameters from the file.

After this point, the network is ready to be used for prediction on new data. For example, the test set can be passed through the network and the accuracy of the model can be calculated on the test set.

```

1. dataiter = iter(testloader)
2. images, labels = next(dataiter)
3.
4.

```

```

5. # print images
6. imshow(torchvision.utils.make_grid(images[:4]), stats)
7. print('GroundTruth: ', ' '.join(f'{classes[labels[j]]:5s}' for j in range(4)))
8. outputs = net(images)

```

This code loads a batch of test data from the testloader (a PyTorch DataLoader that was previously defined and used to load the test data), and assigns the first batch of images and labels to the variables images and labels.

The function imshow(torchvision.utils.make\_grid(images[:4]), stats) is used to display the first 4 images in the batch using the torchvision.utils.make\_grid function to create a grid of images and the imshow function to display the grid.

The line print('GroundTruth: ', ' '.join(f'{classes[labels[j]]:5s}' for j in range(4))) prints the true labels of the first 4 images in the batch. classes is a list of class names and labels[j] is the index of the true class for the j-th image.

The final line outputs = net(images) applies the neural network to the batch of images and assigns the output of the network to the variable outputs. The outputs will be a tensor of shape (batch\_size, number of classes) representing the predicted class scores for each image in the batch.

```

1. _, predicted = torch.max(outputs, 1)
2.
3. print('Predicted: ', ' '.join(f'{classes[predicted[j]]:5s}'
4.                               for j in range(4)))

```

This code performs a prediction on the batch of test images using the loaded neural network.

The torch.max(outputs, 1) function is used to find the class with the highest predicted score for each image in the batch. The first element of the returned tuple is the maximum values and the second element is the indices of the maximum values along the specified dimension (in this case, dimension 1, which corresponds to the classes)

The predicted variable contains the class indices of the highest predicted scores. These predicted class indices are then used to print the predicted class names of the first 4 images in the batch using the classes list, which maps class indices to class names.

It is possible to calculate the accuracy of the model by comparing the predicted labels with the true labels and counting the number of correct predictions.

```

1. correct = 0
2. total = 0
3. # since we're not training, we don't need to calculate the gradients for our outputs
4. with torch.no_grad():
5.     for data in testloader:
6.         images, labels = data
7.         # calculate outputs by running images through the network
8.         outputs = net(images)

```



```

9.         # the class with the highest energy is what we choose as prediction
10.        _, predicted = torch.max(outputs.data, 1)
11.        total += labels.size(0)
12.        correct += (predicted == labels).sum().item()
13.
14.    print(f'Accuracy of the network on the 10000 test images: {100 * correct / total} %')

```

This code calculates the accuracy of the neural network on the test dataset.

The accuracy is calculated by comparing the predicted labels, obtained by finding the class with the highest predicted score for each image in the batch using `torch.max(outputs.data, 1)`, with the true labels.

The outer loop iterates over the test data using the testloader PyTorch DataLoader. For each batch of test data, the network is run on the images to obtain predicted scores, the predicted class is obtained by taking the argmax of the predicted scores and the predicted class is compared to the true labels.

The number of correct predictions is accumulated in the correct variable, and the total number of predictions is accumulated in the total variable

Finally, the accuracy is calculated by dividing the number of correct predictions by the total number of predictions and printing the result with a message.

It is worth noting that the `torch.no_grad()` context manager is used when evaluating the model on the test set, since we don't need gradients for the forward pass during evaluation, this will reduce the memory usage and speed up the evaluation process.

```

1. # prepare to count predictions for each class
2. correct_pred = {classname: 0 for classname in classes}
3. total_pred = {classname: 0 for classname in classes}
4.
5. # again no gradients needed
6. with torch.no_grad():
7.     for data in testloader:
8.         images, labels = data
9.         outputs = net(images)
10.        _, predictions = torch.max(outputs, 1)
11.        # collect the correct predictions for each class
12.        for label, prediction in zip(labels, predictions):
13.            if label == prediction:
14.                correct_pred[classes[label]] += 1
15.                total_pred[classes[label]] += 1
16.
17.
18. # print accuracy for each class
19. for classname, correct_count in correct_pred.items():
20.     accuracy = 100 * float(correct_count) / total_pred[classname]
21.     print(f'Accuracy for class: {classname:5s} is {accuracy:.1f} %')

```

This code calculates the accuracy of the neural network on the test dataset for each class separately.

It prepares to count the correct and total predictions for each class using a dictionary comprehension.

The outer loop iterates over the test data using the testloader PyTorch DataLoader. For each batch of test data, the network is run on the images to obtain predicted scores, the predicted class is obtained by taking the argmax of the predicted scores. Then it collects the correct predictions for each class by comparing the predicted class to the true label.

Finally, the accuracy for each class is calculated by dividing the number of correct predictions by the total number of predictions for that class, and then printing the result with a message.

This way you can see how well the model is performing on each individual class, which can be useful for identifying potential issues or areas for improvement in the model.